

Monte Carlo Methods

- not assume complete knowledge of the environment, requires no prior knowledge of the environment's dynamics
- Monte Carlo methods require only *experience*—sample sequences of states, actions, and rewards from actual or simulated interaction with an environment.
- applicable only to episodic tasks
- use empirical mean return rather than expected return
- no bootstrapping, estimates for each state are independent

Monte Carlo Prediction

- simply to average the returns observed after visits to that state
- As more returns are observed, the average should converge to the expected value.
- 正如 DP 的预期收益，它是所有情况的加权平均；现在只有一些样例，自然从这些里面取平均

First-visit MC

- given a set of episodes obtained by following π and passing through s
 - s may be visited multiple times in the same episode
 - call the first time it is visited in an episode the first visit to s
 - The first-visit MC method estimates $v_\pi(s)$ as the average of the returns following first visits to s
- every-visit MC method averages the returns following all visits to s

- **First-visit MC prediction, for estimating $V \approx v_\pi$**

Input: a policy π to be evaluated

Initialize:

$V(s) \in \mathbb{R}$, arbitrarily, for all $s \in \mathcal{S}$

$Returns(s) \leftarrow$ an empty list, for all $s \in \mathcal{S}$

Loop forever (for each episode):

Generate an episode following π : $S_0, A_0, R_1, S_1, A_1, R_2, \dots, S_{T-1}, A_{T-1}, R_T$

$G \leftarrow 0$

Loop for each step of episode, $t = T-1, T-2, \dots, 0$:

$G \leftarrow \gamma G + R_{t+1}$

Unless S_t appears in S_0, S_1, \dots, S_{t-1} :

Append G to $Returns(S_t)$

$V(S_t) \leftarrow \text{average}(Returns(S_t))$

- Both first-visit MC and every-visit MC converge to $v_\pi(s)$ as the number of visits (or first visits) to s goes to infinity.

相较于 DP 的优势

- 无需了解环境的 dynamics
 - All of the probabilities must be computed before DP can be applied, and such computations are often complex and error-prone. In contrast, generating the sample games required by Monte Carlo methods is easy.
 - estimate for one state does not build upon the estimate of any other state
- 可能只对某些状态感兴趣，将这些状态作为起点，再计算平均回报

Monte Carlo Estimation of Action Values

- 与 prediction 类似，talk about visits to a state-action pair rather than to a state
- many state-action pairs may never be visited

- To compare alternatives we need to estimate the value of all the actions from each state, not just the one we currently favor. (确定性策略的情况)
- *exploring starts* : the episodes start in a state–action pair, and that every pair has a nonzero probability of being selected as the start
- alternative approach : policies are stochastic with a nonzero probability of selecting all actions in each state
按 exploring starts 的方式探索状态行为对的代替方法, 随机性策略

Monte Carlo Control

- Policy improvement is done by making the policy greedy with respect to the current value function.

$$\begin{aligned}\pi(s) &\doteq \arg \max_a q(s, a) \\ q_{\pi_k}(s, \pi_{k+1}(s)) &= q_{\pi_k}\left(s, \arg \max_a q_{\pi_k}(s, a)\right) \\ &= \max_a q_{\pi_k}(s, a) \\ &\geq q_{\pi_k}(s, \pi_k(s)) \\ &\geq v_{\pi_k}(s)\end{aligned}$$

- Monte Carlo methods can be used to find optimal policies given only sample episodes and no other knowledge of the environment's dynamics.

• policy evaluation operates on an infinite number of episodes

– first approach

- 足够多的 steps 可以保证 policy evaluation 的误差足够小
- 这种方式在面对很简单的问题时, 可能需要过多的 episodes

– second approach

- 在 policy evaluation 未完成的时候就去 policy improvement
- only one iteration of iterative policy evaluation is performed between each step of policy improvement
- alternate between improvement and evaluation steps for single states
- 类似于 value iteration (GPI 的思想), DP用了全部的信息, MC用部分的信息, 更新方式差不多

• Monte Carlo ES

After each episode, the observed returns are used for policy evaluation, and then the policy is improved at all the states visited in the episode.

Monte Carlo ES (Exploring Starts), for estimating $\pi \approx \pi_*$

Initialize:

$\pi(s) \in \mathcal{A}(s)$ (arbitrarily), for all $s \in \mathcal{S}$
 $Q(s, a) \in \mathbb{R}$ (arbitrarily), for all $s \in \mathcal{S}, a \in \mathcal{A}(s)$
 $Returns(s, a) \leftarrow$ empty list, for all $s \in \mathcal{S}, a \in \mathcal{A}(s)$

Loop forever (for each episode):

Choose $S_0 \in \mathcal{S}, A_0 \in \mathcal{A}(S_0)$ randomly such that all pairs have probability > 0
 Generate an episode from S_0, A_0 , following π : $S_0, A_0, R_1, \dots, S_{T-1}, A_{T-1}, R_T$
 $G \leftarrow 0$

Loop for each step of episode, $t = T-1, T-2, \dots, 0$:

$G \leftarrow \gamma G + R_{t+1}$

Unless the pair S_t, A_t appears in $S_0, A_0, S_1, A_1, \dots, S_{t-1}, A_{t-1}$:

Append G to $Returns(S_t, A_t)$

$Q(S_t, A_t) \leftarrow \text{average}(Returns(S_t, A_t))$

$\pi(S_t) \leftarrow \arg \max_a Q(S_t, a)$

- first visit

- 算 Q 平均值, update incrementally 的方式更高效
- all the returns for each state-action pair are accumulated and averaged, irrespective of what policy was in force when they were observed

Monte Carlo Control without Exploring Starts

- In on-policy control methods the policy is generally *soft*, meaning that $\forall s \in \mathcal{S} \ a \in \mathcal{A}(s), \pi(a|s) > 0$, but gradually shifted closer and closer to a deterministic optimal policy.
- ϵ -greedy policies are examples of ϵ -soft policies

On-policy first-visit MC control

On-policy first-visit MC control (for ϵ -soft policies), estimates $\pi \approx \pi_*$

Algorithm parameter: small $\epsilon > 0$

Initialize:

$\pi \leftarrow$ an arbitrary ϵ -soft policy

$Q(s, a) \in \mathbb{R}$ (arbitrarily), for all $s \in \mathcal{S}, a \in \mathcal{A}(s)$

$Returns(s, a) \leftarrow$ empty list, for all $s \in \mathcal{S}, a \in \mathcal{A}(s)$

Repeat forever (for each episode):

Generate an episode following π : $S_0, A_0, R_1, \dots, S_{T-1}, A_{T-1}, R_T$

$G \leftarrow 0$

Loop for each step of episode, $t = T-1, T-2, \dots, 0$:

$G \leftarrow \gamma G + R_{t+1}$

Unless the pair S_t, A_t appears in $S_0, A_0, S_1, A_1, \dots, S_{t-1}, A_{t-1}$:

Append G to $Returns(S_t, A_t)$

$Q(S_t, A_t) \leftarrow \text{average}(Returns(S_t, A_t))$

$A^* \leftarrow \arg\max_a Q(S_t, a)$ (with ties broken arbitrarily)

For all $a \in \mathcal{A}(S_t)$:

$$\pi(a|S_t) \leftarrow \begin{cases} 1 - \epsilon + \epsilon/|\mathcal{A}(S_t)| & \text{if } a = A^* \\ \epsilon/|\mathcal{A}(S_t)| & \text{if } a \neq A^* \end{cases}$$

- 随机序列
- ϵ -greedy policy

$$\begin{aligned} q_\pi(s, \pi'(s)) &= \sum_a \pi'(a|s) q_\pi(s, a) \\ &= \frac{\epsilon}{|\mathcal{A}(s)|} \sum_a q_\pi(s, a) + (1 - \epsilon) \max_a q_\pi(s, a) \\ &\geq \frac{\epsilon}{|\mathcal{A}(s)|} \sum_a q_\pi(s, a) + (1 - \epsilon) \sum_a \frac{\pi(a|s) - \frac{\epsilon}{|\mathcal{A}(s)|}}{1 - \epsilon} q_\pi(s, a) \\ &= \frac{\epsilon}{|\mathcal{A}(s)|} \sum_a q_\pi(s, a) - \frac{\epsilon}{|\mathcal{A}(s)|} \sum_a q_\pi(s, a) + \sum_a \pi(a|s) q_\pi(s, a) \\ &= q_\pi(s, \pi(s)) \end{aligned}$$

在 Q 更新过后, 原策略选出的动作不一定是 Q 值最大的

- 将随机选取动作的形式移到环境内部, 环境以 ϵ 的概率不按输入的动作执行, 等概率随机选取另一个动作; 和原环境收敛得到的值函数是一样的

Off-policy Prediction via Importance Sampling

control method dilemma

- 想要学到最优的行为, 要执行非最优的行为来探索, 以找到最优的行为
- How can they learn about the optimal policy while behaving according to an exploratory policy?

on policy

The on-policy approach in the preceding section is actually a compromise—it learns action values not for the optimal policy, but for a near-optimal policy that still explores.

如果收敛到了最优，就无法探索了，失去了再次提高的可能

– off policy

- the policy being learned about is called the *target policy*
- the policy used to generate behavior is called the *behavior policy*
- An advantage of this separation is that the target policy may be deterministic (e.g., greedy), while the behavior policy can continue to sample all possible actions.
- learning is from data “off” the target policy
- often be applied to learn from data generated by a conventional non-learning controller, or from a human expert
可以从人类经验中学习策略
- Target policy becomes a deterministic optimal policy while the behavior policy remains stochastic and more exploratory.
行为策略可以一直探索，不影响目标策略的收敛

– contrast

- off-policy methods are often of greater variance and are slower to converge
- off-policy methods are more powerful and general
- Off-policy include on-policy methods as the special case in which the target and behavior policies are the same.

• prediction problem

- estimate v_π or q_π , but all we have are episodes following another policy b , where $b \neq \pi$
- both policies are considered fixed and given
- the assumption of *coverage* : require that $\pi(a|s) > 0$ implies $b(a|s) > 0$ (every action taken under π is also taken, at least occasionally, under b)
 π 中遇到的状态-动作对在 b 中一定可以遇到

• importance sampling

- Given a starting state S_t , the probability of the subsequent state-action trajectory, $A_t, S_{t+1}, A_{t+1}, \dots, S_T$, occurring under any policy π is (在某策略下某个序列出现的概率)

$$\begin{aligned} \Pr\{A_t, S_{t+1}, A_{t+1}, \dots, S_T | S_t, A_{t:T-1} \sim \pi\} \\ &= \pi(A_t | S_t) p(S_{t+1} | S_t, A_t) \pi(A_{t+1} | S_{t+1}) \cdots p(S_T | S_{T-1}, A_{T-1}) \\ &= \prod_{k=t}^{T-1} \pi(A_k | S_k) p(S_{k+1} | S_k, A_k) \end{aligned}$$

$p()$ here is the state-transition probability function

- *importance-sampling ratio* : relative probability of their trajectories occurring under the target and behavior policies

$$\rho_{t:T-1} \doteq \frac{\prod_{k=t}^{T-1} \pi(A_k | S_k) p(S_{k+1} | S_k, A_k)}{\prod_{k=t}^{T-1} b(A_k | S_k) p(S_{k+1} | S_k, A_k)} = \prod_{k=t}^{T-1} \frac{\pi(A_k | S_k)}{b(A_k | S_k)}$$

depending only on the two policies and the sequence, not on the $p()$ (MDP)

- $\mathbb{E}[G_t | S_t = s] = v_b(s)$, $\mathbb{E}[\rho_{t:T-1} G_t | S_t = s] = v_\pi(s)$

• Monte Carlo algorithm

- the set of all time steps in which state s is visited, denoted $\mathcal{J}(s)$ (for an every-visit method)
- $T(t)$ denote the first time of termination following time t
- G_t denote the return after t up through $T(t)$
- *ordinary importance sampling*

$$V(s) \doteq \frac{\sum_{t \in \mathcal{J}(s)} \rho_{t:T(t)-1} G_t}{|\mathcal{J}(s)|}$$

- unbiased for first-visit methods, biased for every-visit methods (the bias falls asymptotically to zero as the number of samples increases)
- the variance of ordinary importance sampling is in general unbounded because the variance of the ratios can be unbounded
- *weighted importance sampling*

$$V(s) \doteq \frac{\sum_{t \in \mathcal{J}(s)} \rho_{t:T(t)-1} G_t}{\sum_{t \in \mathcal{J}(s)} \rho_{t:T(t)-1}}$$

- both first-visit methods and every-visit methods are biased
- the bias converges asymptotically to zero as the number of samples increases
- the variance of the estimator converges to zero
- In practice, the weighted estimator usually has dramatically lower variance and is strongly preferred.
- ordinary importance sampling is easier to extend to the approximate methods using function approximation
- In practice, every-visit methods are often preferred because they remove the need to keep track of which states have been visited and because they are much easier to extend to approximations.

Incremental Implementation

$$V_n \doteq \frac{\sum_{k=1}^{n-1} W_k G_k}{\sum_{k=1}^{n-1} W_k}, \quad n \geq 2$$

$$V_{n+1} \doteq V_n + \frac{W_n}{C_n} [G_n - V_n], \quad n \geq 1$$

$$C_{n+1} \doteq C_n + W_{n+1}, C_0 \doteq 0$$

for the off-policy case (on-policy case, W is always 1)

• Off-policy MC prediction

Off-policy MC prediction (policy evaluation) for estimating $Q \approx q_\pi$

Input: an arbitrary target policy π

Initialize, for all $s \in \mathcal{S}$, $a \in \mathcal{A}(s)$:

$Q(s, a) \in \mathbb{R}$ (arbitrarily)

$C(s, a) \leftarrow 0$

Loop forever (for each episode):

$b \leftarrow$ any policy with coverage of π

Generate an episode following b : $S_0, A_0, R_1, \dots, S_{T-1}, A_{T-1}, R_T$

$G \leftarrow 0$

$W \leftarrow 1$

Loop for each step of episode, $t = T-1, T-2, \dots, 0$, while $W \neq 0$:

$G \leftarrow \gamma G + R_{t+1}$

$C(S_t, A_t) \leftarrow C(S_t, A_t) + W$

$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \frac{W}{C(S_t, A_t)} [G - Q(S_t, A_t)]$

$W \leftarrow W \frac{\pi(A_t|S_t)}{b(A_t|S_t)}$

$\rho_{t:T-1}$ 就是连乘的形式

Off-policy Monte Carlo Control

Off-policy MC control, for estimating $\pi \approx \pi_*$

Initialize, for all $s \in \mathcal{S}$, $a \in \mathcal{A}(s)$:

$Q(s, a) \in \mathbb{R}$ (arbitrarily)

$C(s, a) \leftarrow 0$

$\pi(s) \leftarrow \arg\max_a Q(s, a)$ (with ties broken consistently)

Loop forever (for each episode):

$b \leftarrow$ any soft policy

Generate an episode using b : $S_0, A_0, R_1, \dots, S_{T-1}, A_{T-1}, R_T$

$G \leftarrow 0$

$W \leftarrow 1$

Loop for each step of episode, $t = T-1, T-2, \dots, 0$:

$G \leftarrow \gamma G + R_{t+1}$

$C(S_t, A_t) \leftarrow C(S_t, A_t) + W$

$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \frac{W}{C(S_t, A_t)} [G - Q(S_t, A_t)]$

$\pi(S_t) \leftarrow \arg\max_a Q(S_t, a)$ (with ties broken consistently)

If $A_t \neq \pi(S_t)$ then exit inner Loop (proceed to next episode)

$W \leftarrow W \frac{1}{b(A_t|S_t)}$

- based on GPI and weighted importance sampling
- The target policy $\pi \approx \pi^*$ is the greedy policy with respect to Q , which is an estimate of q_π
- The behavior policy b can be anything, but in order to assure convergence of π to the optimal policy, an infinite number of returns must be obtained for each pair of state and action.

出现无限多次保证收敛到最优

- A potential problem is that this method learns only from the tails of episodes, when all of the remaining actions in the episode are greedy. If non-greedy actions are common, then learning will be slow, particularly for states appearing in the early portions of long episodes.

Discounting-aware Importance Sampling

- reduce the variance of off-policy estimators
- *flat partial returns*: $\bar{G}_{t:h} \doteq R_{t+1} + R_{t+2} + \dots + R_h$, $0 \leq t < h \leq T$
 - “flat” denotes the absence of discounting
 - “partial” denotes that these returns do not extend all the way to termination but instead stop at h , called the *horizon*
 - T is the time of termination of the episode

$$G_t \doteq R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \dots + \gamma^{T-t-1} R_T$$

$$\begin{aligned} &= (1 - \gamma) R_{t+1} \\ &\quad + (1 - \gamma) \gamma (R_{t+1} + R_{t+2}) \\ &\quad + (1 - \gamma) \gamma^2 (R_{t+1} + R_{t+2} + R_{t+3}) \\ &\quad \dots \\ &\quad + (1 - \gamma) \gamma^{T-t-2} (R_{t+1} + R_{t+2} + \dots + R_{T-1}) \\ &\quad + \gamma^{T-t-1} (R_{t+1} + R_{t+2} + \dots + R_T) \end{aligned}$$

$$= (1 - \gamma) \sum_{h=t+1}^{T-1} \gamma^{h-t-1} \bar{G}_{t:h} + \gamma^{T-t-1} \bar{G}_{t:T}$$

- discounting-aware importance sampling estimators (ordinary and weighted)

$$V(s) \doteq \frac{\sum_{t \in \mathcal{J}(s)} \left((1 - \gamma) \sum_{h=t+1}^{T(t)-1} \gamma^{h-t-1} \rho_{t:h-1} \bar{G}_{t:h} + \gamma^{T(t)-t-1} \rho_{t:T(t)-1} \bar{G}_{t:T(t)} \right)}{|\mathcal{J}(s)|}$$

$$V(s) \doteq \frac{\sum_{t \in \mathcal{J}(s)} \left((1 - \gamma) \sum_{h=t+1}^{T(t)-1} \gamma^{h-t-1} \rho_{t:h-1} \bar{G}_{t:h} + \gamma^{T(t)-t-1} \rho_{t:T(t)-1} \bar{G}_{t:T(t)} \right)}{\sum_{t \in \mathcal{J}(s)} \left((1 - \gamma) \sum_{h=t+1}^{T(t)-1} \gamma^{h-t-1} \rho_{t:h-1} + \gamma^{T(t)-t-1} \rho_{t:T(t)-1} \right)}$$

- 优化了 $\gamma = 0$ 的情况 (详见书)

Per-decision Importance Sampling

$$\begin{aligned}\rho_{t:T-1}G_t &= \rho_{t:T-1}(R_{t+1} + \gamma R_{t+2} + \cdots + \gamma^{T-t-1}R_T) \\ &= \rho_{t:T-1}R_{t+1} + \gamma\rho_{t:T-1}R_{t+2} + \cdots + \gamma^{T-t-1}\rho_{t:T-1}R_T\end{aligned}$$

计算 $\rho_{t:T-1}R_{t+1}$ 时, $\frac{\pi(A_k|S_k)}{b(A_k|S_k)}$ ($t+1 \leq k \leq T-1$) 其实是无关的, 不需要知道后面的情况

$$\mathbb{E}\left[\frac{\pi(A_k|S_k)}{b(A_k|S_k)}\right] \doteq \sum_a b(a|S_k) \frac{\pi(a|S_k)}{b(a|S_k)} = \sum_a \pi(a|S_k) = 1$$

$$\mathbb{E}[\rho_{t:T-1}R_{t+k}] = \mathbb{E}[\rho_{t:t+k-1}R_{t+k}]$$

$$\mathbb{E}[\rho_{t:T-1}G_t] = \mathbb{E}[\tilde{G}_t]$$

$$\tilde{G}_t = \rho_{t:t}R_{t+1} + \gamma\rho_{t:t+1}R_{t+2} + \gamma^2\rho_{t:t+2}R_{t+3} + \cdots + \gamma^{T-t-1}\rho_{t:T-1}R_T$$

$$V(s) \doteq \frac{\sum_{t \in \mathcal{T}(s)} \tilde{G}_t}{|\mathcal{T}(s)|}$$

Summary

- Monte Carlo methods provide an alternative policy evaluation process. Rather than use a model to compute the value of each state, they simply average many returns that start in the state.
- Despite their conceptual simplicity, off-policy Monte Carlo methods for both prediction and control remain unsettled and are a subject of ongoing research.